

DDA3020 Machine Learning: Lecture 1 Introduction

Baoyuan Wu
School of Artificial Intelligence, CUHK-SZ

January 05/07, 2026

Outline

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning

Instructors and teaching assistants (Session 1)

Time and venue:

- Lecture: Monday/Wednesday 9:00–10:20 am, Administration Bldg E101
- Tutorial: **TBD**

Instructor:

- **Baoyuan Wu** (wubaoyuan@cuhk.edu.cn)
- Personal homepage: <https://sites.google.com/site/baoyuanwu2015/home>
- Office hour (OH): 10:30–11:30 am, Monday in the even weeks (2/4/6/8/10/12), TA 412a

TAs:

- **Meixi Zheng**: 223040254@link.cuhk.edu.cn
OH: Mon 11:00-12:00, TXA 101
- **Yuan Xie**: 224040374@link.cuhk.edu.cn
OH: Tue 16:00-17:00, DaoYuan 318
- **Jun'an Chen**: 225040494@link.cuhk.edu.cn
OH: Tue 16:00-17:00, DaoYuan 318
- **Yueci Deng**: 225040502@link.cuhk.edu.cn
OH: Tue 14:00-15:00, TXB 609
- **Dekun Chen**: dekunchen@link.cuhk.edu.cn
OH: Wed 14:00-15:00, TXC 604
- **Fanzeng Xia**: 223040232@link.cuhk.edu.cn
OH: Mon 9:00-10:00, TXB 611

Instructors and teaching assistants (Session 2)

Time and venue:

- Lecture: Wednesday/Friday 10:30–11:50 am, Administration Bldg E101
- Tutorial: **TBD**
- **Kui Jia** (kuijia@cuhk.edu.cn)
- Personal homepage: <http://kuijia.site/>
- Office hour (OH): 14:00–15:00 pm, Wednesday in even weeks (2/4/6/8/10/12), TxC 607

TAs:

- **Meixi Zheng**: 223040254@link.cuhk.edu.cn
OH: Mon 11:00-12:00, TXA 101
- **Yuan Xie**: 224040374@link.cuhk.edu.cn
OH: Tue 16:00-17:00, DaoYuan 318
- **Jun'an Chen**: 225040494@link.cuhk.edu.cn
OH: Tue 16:00-17:00, DaoYuan 318
- **Yueci Deng**: 225040502@link.cuhk.edu.cn
OH: Tue 14:00-15:00, TXB 609
- **Dekun Chen**: dekunchen@link.cuhk.edu.cn
OH: Wed 14:00-15:00, TxC 604
- **Fanzeng Xia**: 223040232@link.cuhk.edu.cn
OH: Mon 9:00-10:00, TXB 611

Agenda with tutorials

Week	Content	Homework	Tutorial
W1	Introduction		None
W2	Review of Probability & Linear Algebra etc.		python and sklearn
W3	Linear Regression I		Linear regression coding
W4	Linear Regression II & Logistic Regression	H1 release	Logistic regression coding
W5	Support Vector Machines		SVM coding
W6	Decision Tree and Random Forest		H1
W7	Neural Networks I (MLP & CNN)		Pytorch coding
W8	Neural Networks II (RNN & Transformer)	H2 release	Pytorch coding
W9	Over-fitting, Bias-Variance Trade-off		Demo
W10	Performance Evaluation		H2
W11	Intro. to Unsupervised Learning, K-Means		K-means coding
W12	Mixture Models, EM algorithm	H3 release	GMM coding
W13	PCA		PCA coding
W14	Review		

Textbooks:

- Required: Andriy Burkov, “The Hundred-Page Machine Learning Book”, 2019. (read first, buy later: [Link](#))
- Required: K. Murphy. Machine Learning: A Probabilistic Perspective. MIT Press, 2012. ([Link](#))
- Recommended: Andreas C. Muller and Sarah Guido, “Introduction to Machine Learning with Python: A Guide for Data Scientists”, O’Reilly Media, Inc., 2017. (The PDF will be uploaded to the BB system, only used for learning)
- Recommended: Jeff Leek, “The Elements of Data Analytic Style: A guide for people who want to analyze data”, Lean Publishing, 2015. ([Link](#))
- Recommended: C. Bishop. Pattern Recognition and Machine Learning. Springer, 2011. ([Link](#))

Grading policy

Grading:

- 25% Written homework assignments (3 times)
- 25% Programming homework assignments (Python/Scikit-learn) (3 times)
- 50% Final exam

Submission: You have 2 weeks to **independently** complete each assignment (written + programming). Late submission will get discounted score: $(0, 48]$ hours $\rightarrow$ 50%; $(48, \infty)$ hours $\rightarrow$ 0%.

Requirements to finish the homework smoothly: basic knowledge of probability and linear algebra, familiarity with Python programming, and correct understanding of the teaching content.

Plagiarism: Zero mark is given for the whole assignment (including written and programming) in the first plagiarism case. Students will **FAIL** the whole course for repeated plagiarism. **Note:** If there are heavy overlaps between two answers, then both will be identified as plagiarism (we don't have time to distinguish). Thus, discussions are encouraged, but you must finish the assignment by yourself, and don't share your answer with others.


The use of LLMs or AI agents to **assist** your learning, homework, or programming is **permitted**.

However, you must clearly **declare** such usage. Any work identified as AI-generated without a proper declaration will be treated as **plagiarism**.


Be the master of AI, not its slave!

Students' background

- The students come from **diverse years and majors**, thus, may have different expectations and experiences.



Students' background



Course difficulty

- Machine learning is a **complex** and **evolving** discipline. Our course will cover the **fundamental concepts and algorithms** in machine learning.
- **Fundamental $\neq$ SIMPLE**. Machine learning is built on several other foundation subjects, including **probability, linear algebra, optimization, computer science**, *etc.*
- If you find it **difficult** to understand the teaching content, please don't hesitate to tell me or TAs, and we will help you overcome the challenges. If you feel the content is **too easy**, no worry, we can provide further reading materials on advanced and emerging topics (*e.g.*, deep learning, adversarial machine learning). Wish each of you an enjoyable learning experience!

Smooth communications between you and us are very important to build a successful course. If any questions/difficulties/suggestions, you are welcome to

- Talk to me directly, after class or during office hours
- Email or talk to your TAs, questions will be collected and sent to me weekly for responses
- Questions about programming are handled by TAs
- You may also send me emails if you need my help

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning

Definition of machine learning

- Arthur Samuel: “the field of study that gives computers the ability to learn without being explicitly programmed.”
- Tom Mitchell: “A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E.”

Definition of machine learning

Tom Mitchell: “A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E.”

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task T in this setting?

- Classifying emails as spam or not spam.
- Watching you label emails as spam or not spam.
- The number (or fraction) of emails correctly classified as spam/not spam.
- None of the above – this is not a machine learning problem.

ML is a branch of artificial intelligence

- **Artificial intelligence (AI)** is intelligence demonstrated by **machines** (*e.g.*, computer, robots), unlike the natural intelligence displayed by humans and animals, which involves consciousness and emotionality.
- AI covers many topics, such as machine learning (ML), computer vision (CV), natural language processing (NLP), and speech processing.




Figure: Machine learning is one of the most important branches of artificial intelligence.

Connection with other disciplines




Figure: Machine learning is interdisciplinary.

Applications of machine learning




Figure: Machine learning has been widely used in many mission-critical tasks, such as Game (AlphaGo), protein structure prediction (AlphaFold2), EEG signal processing, medical image diagnosis, face recognition, etc.


Two basic paradigms of ML

Given some data $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \dots$, you can do

- **Supervised learning:** you are also provided some human-labeled outputs $\mathbf{y}_1, \mathbf{y}_2, \mathbf{y}_3, \dots$, and your task is to learn a mapping function from one input data $\mathbf{x}_i$ to one output $\mathbf{y}_i$. **Learning from teacher**
- **Unsupervised learning:** your task is to build/learn a good model of $\mathbf{x}$, such that some characteristics of the data could be revealed, such as clustering, dimensionality reduction, *etc.* **Learning by oneself**


Other learning paradigms

- **Reinforcement learning:** you can make some actions a_i to change the data $\mathbf{x}_i$ (the state of the environment), and you will receive some rewards/punishments r_i . Gradually, you will automatically learn to make suitable actions for different data to get more rewards. It is one of three basic machine learning paradigms, alongside supervised learning and unsupervised learning. **Learning from rewards/punishments**



In this course, we will focus on **supervised learning** and **unsupervised learning**.

Supervised vs. unsupervised learning




	Supervised Learning (with labels)	Unsupervised Learning (without labels)
Discrete	Classification	Clustering
Continuous	Regression	Dimensionality Reduction

Refer to:

<https://towardsdatascience.com/supervised-vs-unsupervised-learning-14f68e32ea8d>

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning**
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning

Supervised learning




Supervised learning

In **supervised learning**, the **dataset** is the collection of **labeled examples**, denoted as $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$,

$$\mathbf{x}_i = [x_i^{(1)}, \dots, x_i^{(j)}, \dots, x_i^{(D)}]^\top, \quad i = 1, \dots, N$$

- Each element $\mathbf{x}_i$ is called a **feature vector**: it is a vector in which each dimension $j = 1, \dots, D$ contains a value that describes the example somehow.
- The **label** y_i can be either an element belonging to a **finite** set of classes $\{1, 2, \dots, C\}$, or a **real number**.


How supervised learning works



General procedure of supervised learning:


- **Data collection:** $\{(x_i, y_i)\}_{i=1}^N$
- **Training:** conducting model training on the training data to learn the model parameters
- **Inference (test):** Using the trained model to predict the output of unseen data x

Regression example



- **Task:** To predict the price of diamond as a function of mass
- **Performance:** Accuracy of predicted values
- **Experience:** Historical data

Classification example



- **Task:** To classify the input data into two categories
- **Performance:** Classification accuracy
- **Experience:** Historical data

Regression vs. classification

Suppose that you are running a company, and you want to develop learning algorithms to address the following two problems.

Problem 1: You have a large inventory of identical items. You want to predict how many of these items will sell over the next 3 months.

Problem 2: You'd like a program to examine individual customer accounts, and for each account, decide if it has been hacked or compromised.


Should you treat these as classification or regression problems?

- Treat both as classification problems.
- Treat problem 1 as classification, and problem 2 as regression.
- Treat problem 1 as regression, and problem 2 as classification.
- Treat both as regression problems.

Regression with **continuous output** vs. Classification with **finite and discrete or categorical outputs**

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning**
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning


Unsupervised learning



Unsupervised learning

- In **unsupervised learning**, the **dataset** is a collection of **unlabeled examples**, *i.e.*, $\{\mathbf{x}_i\}_{i=1}^N$.
- Again, $\mathbf{x}$ is a feature vector, and the goal of **an unsupervised learning algorithm** is to create a **model** that takes a feature vector $\mathbf{x}$ as input and **either transforms it into another vector or into a value** that can be used to solve a practical problem.
- Its **main task** is to **analyze the structure of data** for future inference.

How unsupervised learning works




Clustering example



- **Task:** to partition a set of unlabeled points to clusters.
- **Performance (for test data):**
 - points within the same cluster are **close** to each other
 - points from different clusters are **far** from each other
 - the clusters have an appropriate coverage of all data
- **Experience:** available data

The key question is how to define and measure **close** or **far**, depending on what features are adopted (*e.g.*, global Euclidean distance, local distance, shape ...)

Dimensionality reduction example



The purposes of dimensionality reduction:

- **Data simplification:** non-linear $\rightarrow$ linear
- **Data visualization:** high dimensional $\rightarrow$ low-dimensional
- **Reduce noise:** some dimensions of the input data may be noises
- **Variable selection for prediction:** learn a sparse model, if there are redundancies among different dimensions

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning**
- 6 Practice for machine learning

Basic concepts in supervised learning

- **Data for supervised learning:**

- **Training set:** $D_{train} = \{(\mathbf{x}_i, y_i)\}_i^n$, with $\mathbf{x}_i \in \mathcal{X}$ being the feature presentation (could be scalar, image/video, text sequence, graph, cloud point, *etc.*), $y_i \in \mathcal{Y}$ being the supervision/ground-truth value (could be a discrete label, continuous value, vector, sequence, tensor, *etc.*).
- **Testing set:** $D_{test} = \{(\mathbf{x}_i, y_i)\}_i^m$ is used to evaluate the performance of the trained model.

- **Data for unsupervised learning**

- **Training set:** $D_{train} = \{\mathbf{x}_i\}_i^n$, with $\mathbf{x}_i \in \mathcal{X}$ being the feature presentation (could be scalar, image/video, text sequence, graph, cloud point, *etc.*).
- **Testing set:** $D_{test} = \{\mathbf{x}_i\}_i^m$ is used to evaluate the performance of the trained model.

- **Independent and identically distributed (i.i.d.) assumption:** In standard machine learning, it assumes that all samples are observations (or realizations) of [independent and identical random variables](#), and training and testing sets follow [the same distribution](#).


Basic concepts in supervised learning

- **Target function** $t : \mathcal{X} \rightarrow \mathcal{Y}$: the ground-truth mapping function from the input to the output behind the training/testing data. It is unknown, and our goal is to find it.
- **Hypothesis** h : A hypothesis is a candidate function that describes the unknown target function, for example

$$h(\mathbf{x}) = 1 \times x_1 + 2 \times x_2 = [1, 2] \cdot \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- **Hypothesis space** $\mathcal{H}$: Hypothesis space is the set of all the possible legal hypotheses, for example

$$\mathcal{H}_{\mathbf{w}}(\mathbf{x}) = w_1 \times x_1 + w_2 \times x_2 = \mathbf{w}^\top \mathbf{x}$$



Basic concepts in supervised learning

- **Cost function** is a measure of how good/bad the hypothesis is in terms of its ability to estimate the relationship between $\mathbf{x}$ and y , such as **the square loss** $(h(\mathbf{x}) - y)^2$.
- **Objective function** is the function that we want to optimize (minimize, maximize, or minimax). When we are **minimizing** it, we may also call it the **cost function**, loss function, or error function. In this course, we use these terms interchangeably.
- **Training/learning** is the process of searching a good hypothesis h in the hypothesis space $\mathcal{H}$, through optimizing the objective function on a training set D_{train} using the optimization method, such as

$$h^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{(\mathbf{x}_i, y_i) \in D_{train}} (h(\mathbf{x}_i) - y_i)^2$$

- **Testing/evaluation:** evaluating the performance of the learned model h^* on a test set D_{test} , such as

$$\frac{1}{m} \sum_{(\mathbf{x}_i, y_i) \in D_{test}} (h^*(\mathbf{x}_i) - y_i)^2$$

A general machine learning workflow

- ❶ **Collecting data:** Be it the raw data from Excel, access, text files, etc., this step (gathering past data) forms the foundation of future learning. The better the variety, density, and volume of relevant data, the better the learning prospects for the machine become.
- ❷ **Preprocessing data:** One needs to spend time determining the quality of data and then taking steps for fixing issues such as missing data and treatment of outliers.
- ❸ **Determining the hypothesis space, objective function, optimization method.**
- ❹ **Training:** learning the parameters of the hypothesis function through optimizing the objective function.
- ❺ **Testing:** evaluating the learned model on testing data.
- ❻ **Improving the performance:** This step might involve choosing a different model altogether or introducing more variables to augment the efficiency. That's why a significant amount of time needs to be spent on data collection and preparation.

Review of this lecture

- **Machine learning definition:** “A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E.”
- **Three main paradigms of machine learning:**
 - **Supervised learning:** learning from teacher
 - **Unsupervised learning:** learning by oneself
 - Reinforcement learning: learning from rewards/punishments
 - Many other learning paradigms ...
- **Supervised learning:** regression and classification.
- **Unsupervised learning:** clustering and dimensionality reduction.
- **Some basic concepts**

Other learning paradigms:

- **Reinforcement learning:** an area of machine learning concerned with how intelligent agents ought to take actions in an environment in order to maximize the notion of cumulative reward. ([Link](#))
- **Semi-supervised learning:** labeled data + unlabeled data ([Link](#))
- **Ensemble learning:** learning with multiple ML models, and the final prediction is obtained by combining these models. **Improving the performance of individual models** 三个臭皮匠，顶个诸葛亮 ([Link](#))
- **Transfer learning:** source domain data + target domain data, useful especially when the target data is insufficient ([Link](#))
- **Federated learning:** learning the model at local servers using local data, then uploading locally updated parameters to the central server to obtain the unified parameters. **Protecting users' privacy** ([Link](#))
- **Machine unlearning:** erasing the effect of some particular training samples from a trained model. **Protecting users' privacy** ([Link](#))

- 1 About this course
- 2 What is machine learning
- 3 Supervised learning
- 4 Unsupervised learning
- 5 Some basic concepts in machine learning
- 6 Practice for machine learning**

Practice: Scikit-Learn

<https://scikit-learn.org/stable/>



Install User Guide API Examples Community More ▾

Go

scikit-learn

Machine Learning in Python

Getting Started

Release Highlights for 1.4

GitHub


- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: Gradient boosting, nearest neighbors, random forest, logistic regression, and more...




Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: Gradient boosting, nearest neighbors, random forest, ridge, and more...




Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, HDBSCAN, hierarchical clustering, and more...



Examples

Dimensionality reduction

Model selection

Preprocessing

Practice: UCI Data Repository

<https://archive.ics.uci.edu/>

UC Irvine
Machine Learning
Repository

Datasets Contribute Dataset About Us

Search datasets...   Login

Welcome to the UC Irvine Machine Learning Repository

We currently maintain 664 datasets as a service to the machine learning community. Here, you can donate and find datasets used by millions of people all around the world!

[VIEW DATASETS](#) [CONTRIBUTE A DATASET](#)

Popular Datasets

**Iris**
A small classic dataset from Fisher, 1936. One of the earliest known datasets used for...
 Classification  150 Instances  4 Features

**Dry Bean**
Images of 13,611 grains of 7 different registered dry beans were taken with a high-res...
 Classification  13.61K Instances  16 Features

**Heart Disease**
4 databases: Cleveland, Hungary, Switzerland, and the VA Long Beach
 Classification  303 Instances  13 Features

**Rice (Cammeo and Osmancik)**
A total of 3810 rice grain's images were taken for the two species, processed and feat...
 Classification  3.81K Instances  7 Features

Adult

New Datasets

**PhiSUIIL Phishing URL (Website)**
PhiSUIIL Phishing URL Dataset is a substantial dataset comprising 134,850 legitimate ...
 Classification  235.8K Instances  54 Features

**RT-IoT2022**
The RT-IoT2022, a proprietary dataset derived from a real-time IoT infrastructure, is in...
 Classification, Regres...  123.12K Instances  84 Features

**Regensburg Pediatric Appendicitis**
This repository holds the data from a cohort of pediatric patients with suspected app...
 Classification  782 Instances  59 Features

**National Poll on Healthy Aging (NPHA)**
This is a subset of the NPHA dataset filtered down to develop and validate machine le...
 Classification  714 Instances  15 Features

Infrared Thermography Temperature

Practice: Kaggle

<https://www.kaggle.com/>

kaggle

[Competitions](#) [Datasets](#) [Models](#) [Code](#) [Discussions](#) [Courses](#) [...](#)

 Search

[Sign In](#)

[Register](#)

Level up with the largest AI & ML community

Join over 17M+ machine learners to share, stress test, and stay up-to-date on all the latest ML techniques and technologies. Discover a huge repository of community-published models, data & code for your next project.

 Register with Google

[Register with Email](#)



Who's on Kaggle?

Learners

Dive into Kaggle courses, competitions & forums.



Developers

Leverage Kaggle's models, notebooks & datasets.



Researchers

Advance ML with our pre-trained model hub & competitions.

